

September 20

1 Advanced Topics in AI Research

Explored cutting-edge developments in artificial intelligence and their implications.

1.1 Reinforcement Learning

The agent-environment interaction follows the Markov Decision Process:

Bellman Equation:

$$V^\pi(s) = \mathbb{E}[R_{t+1} + \gamma V^\pi(S_{t+1}) | S_t = s]$$

Q-Learning Update:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

1.2 Policy Gradient Methods

- **REINFORCE:** $\nabla_\theta J(\theta) = \mathbb{E}[\nabla_\theta \log \pi_\theta(a|s) Q^\pi(s, a)]$
- **Actor-Critic:** Combines value function approximation with policy gradients
- **PPO:** Proximal Policy Optimization with clipped surrogate objective
- **TRPO:** Trust Region Policy Optimization for stable updates

1.3 Meta-Learning (Learning to Learn)

Key approaches:

1. Model-Agnostic Meta-Learning (MAML):

$$\theta' = \theta - \alpha \nabla_\theta L_{T_i}(f_\theta)$$

$$\theta \leftarrow \theta - \beta \nabla_\theta \sum_{T_i} L_{T_i}(f_{\theta'})$$

2. Prototypical Networks: Learn embeddings where classes form clusters

3. Memory-Augmented Networks: External memory for few-shot learning

1.4 Interpretability and Explainable AI

- **LIME:** Local Interpretable Model-agnostic Explanations
- **SHAP:** SHapley Additive exPlanations using game theory
- **Grad-CAM:** Gradient-weighted Class Activation Mapping for CNNs
- **Attention Visualization:** Understanding what transformers focus on

1.5 Future Research Directions

- Causal inference in machine learning
- Federated learning for privacy-preserving AI
- Neuromorphic computing and spiking neural networks
- AI safety and alignment research